

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.simplefilter('ignore')
```

```
In [2]: df = pd.read_excel('Data_Train.xlsx')
```

```
In [3]: df.head()
```

Out[3]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total
0	IndiGo	24/03/2019	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	2h 50m	n
1	Air India	1/05/2019	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	
2	Jet Airways	9/06/2019	Delhi	Cochin	DEL → LKO → BOM → COK	09:25	04:25 10 Jun	19h	
3	IndiGo	12/05/2019	Kolkata	Banglore	CCU → NAG → BLR	18:05	23:30	5h 25m	
4	IndiGo	01/03/2019	Banglore	New Delhi	BLR → NAG → DEL	16:50	21:35	4h 45m	

```
In [4]: df.shape
```

Out[4]: (10683, 11)

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline                10683 non-null  object
1   Date_of_Journey        10683 non-null  object
2   Source                 10683 non-null  object
3   Destination            10683 non-null  object
4   Route                  10682 non-null  object
5   Dep_Time               10683 non-null  object
6   Arrival_Time           10683 non-null  object
7   Duration               10683 non-null  object
8   Total_Stops            10682 non-null  object
9   Additional_Info        10683 non-null  object
10  Price                  10683 non-null  int64
dtypes: int64(1), object(10)
memory usage: 918.2+ KB
```

Missing values

```
In [6]: df.isna().sum()
```

```
Out[6]: Airline                0
Date_of_Journey            0
Source                     0
Destination                 0
Route                       1
Dep_Time                   0
Arrival_Time               0
Duration                   0
Total_Stops                 1
Additional_Info             0
Price                      0
dtype: int64
```

```
In [7]: a = df.isna().any()
na_col = a[a].index
na_col
```

```
Out[7]: Index(['Route', 'Total_Stops'], dtype='object')
```

```
In [8]: df.dropna(inplace = True)
```

```
In [9]: df.duplicated().sum()
```

```
Out[9]: 220
```

There are 220 rows which are duplicates.


```
In [10]: df.drop_duplicates(inplace=True)
```

```
In [11]: print(df['Airline'].value_counts())  
df['Airline'].value_counts().plot(kind = 'bar')
```

Jet Airways	3700
IndiGo	2043
Air India	1694
Multiple carriers	1196
SpiceJet	815
Vistara	478
Air Asia	319
GoAir	194
Multiple carriers Premium economy	13
Jet Airways Business	6
Vistara Premium economy	3
Trujet	1

Name: Airline, dtype: int64

```
Out[11]: <AxesSubplot:>
```



jet Airways , indiGO , Air Inida are the top airlines which passengers prefer.

```
In [12]: airlines_with_very_less_data= (df['Airline'].value_counts() < 14)
1 = airlines_with_very_less_data[airlines_with_very_less_data].index
1
```

```
Out[12]: Index(['Multiple carriers Premium economy', 'Jet Airways Business',
               'Vistara Premium economy', 'Trujet'],
              dtype='object')
```

these are the airlines which data are very less compare to other airlines. Multiple carriers Premium economy , Jet Airways Business , Vistara Premium economy , Trujet

```
In [13]: #converting string dtype to datetime
df['Date_of_Journey'] = pd.to_datetime(df['Date_of_Journey'])
```

```
In [14]: df['Date_of_Journey'].dt.year.value_counts()
```


```
Out[14]: 2019      10462
         Name: Date_of_Journey, dtype: int64
```

we have only 2019 data

```
In [15]: print(df['Source'].value_counts())
df['Source'].value_counts().plot(kind = 'bar')
```

```
Delhi      4345
Kolkata    2860
Banglore   2179
Mumbai      697
Chennai     381
         Name: Source, dtype: int64
```


```
Out[15]: <AxesSubplot:>
```



```
In [16]: print(df['Destination'].value_counts())  
df['Destination'].value_counts().plot(kind = 'bar')
```

```
Cochin      4345  
Banglore    2860  
Delhi       1265  
New Delhi   914  
Hyderabad   697  
Kolkata     381  
Name: Destination, dtype: int64
```


Out[16]: <AxesSubplot:>



From the above bar graph we have to destination delhi , new delhi. but in source only delhi is there so we will replace NewDelhi to Delhi.

```
In [17]: df['Destination'] = df['Destination'].replace('New Delhi' , 'Delhi')
df['Destination'].value_counts().plot(kind = 'bar')
```

Out[17]: <AxesSubplot:>



```
In [18]: df['Route'].value_counts()
```

```
Out[18]: DEL → BOM → COK          2376
BLR → DEL          1536
CCU → BOM → BLR     979
CCU → BLR           724
BOM → HYD           621
...
CCU → VTZ → BLR      1
CCU → IXZ → MAA → BLR 1
BOM → COK → MAA → HYD 1
BOM → CCU → HYD       1
BOM → BBI → HYD       1
Name: Route, Length: 128, dtype: int64
```

There are total 128 different routes.

```
In [19]: print(df['Total_Stops'].value_counts())
```


```
1 stop          5625
non-stop        3475
2 stops         1318
3 stops          43
4 stops          1
Name: Total_Stops, dtype: int64
```

there are total 5 kind of stops but we have only one data for 4 stops so remove that.

```
In [20]: df = df[df['Total_Stops']!='4 stops']
```

```
In [21]: df['Total_Stops'].value_counts().plot(kind = 'bar')
```


Out[21]: <AxesSubplot:>



Replacing 1stops = 1 , 2stops= 2 , 3stops = 3 , non-stops = 0


```
In [22]: df['Total_Stops'] = df['Total_Stops'].str[0].replace('n' , 0).astype('int32')
df['Total_Stops'].value_counts().plot(kind = 'bar')
```

Out[22]: <AxesSubplot:>



```
In [23]: df['Additional_Info'].value_counts().plot(kind = 'bar')
```

Out[23]: <AxesSubplot:>



```
In [24]: df['Dep_Time' ][0:5]
```

Out[24]:

0	22:20
1	05:50
2	09:25
3	18:05
4	16:50

Name: Dep_Time, dtype: object

```
In [25]: # Extracting departure hours and departure minutes
df['dep_hour' ] = df['Dep_Time'].str.split(':', expand = True)[0].astype('int32')
df['dep_min' ] = df['Dep_Time'].str.split(':', expand = True)[1].astype('int32')
```



```
In [26]: df['Arrival_Time'][0:5]
```

```
Out[26]: 0    01:10 22 Mar
1         13:15
2    04:25 10 Jun
3         23:30
4         21:35
Name: Arrival_Time, dtype: object
```

```
In [27]: # Extracting arrival hours and arrival minutes
```

```
df['arrival_hour'] = df['Arrival_Time'].str[0:5].str.split(':', expand = True)[0]
df['arrival_min'] = df['Arrival_Time'].str[0:5].str.split(':', expand = True)[1]
```

```
In [28]: # Extracting Duration hours and duration minutes
```

```
df['Duration_hour'] = df['Duration'].str.replace('h', '').str.replace('m', '').str[0:2]
df['Duration_min'] = df['Duration'].str.replace('h', '').str.replace('m', '').str[3:5]
```

```
In [29]: #calculation total duration hours
```


```
df['total_hours'] = df['Duration_hour'] + df['Duration_min'] / 60
```

EDA


```
In [30]: df.columns
```

```
Out[30]: Index(['Airline', 'Date_of_Journey', 'Source', 'Destination', 'Route',
               'Dep_Time', 'Arrival_Time', 'Duration', 'Total_Stops',
               'Additional_Info', 'Price', 'dep_hour', 'dep_min', 'arrival_hour',
               'arrival_min', 'Duration_hour', 'Duration_min', 'total_hours'],
              dtype='object')
```

```
In [31]: sns.barplot(x = 'Airline' , y = 'Price' , data = df )  
plt.xticks(rotation = 65)  
plt.show()
```




```
In [32]: sns.boxplot(x = 'Airline' , y = 'Price' , data = df )
plt.xticks(rotation = 65)
plt.show()
```




From graph we can see that Jet Airways Business have the highest Price., Apart from the Jet Airways Business almost all are having similar median

```
In [33]: sns.boxplot(x = 'Source' , y = 'Price' , data = df )
plt.xticks(rotation = 65)
plt.show()
```



```
In [34]: sns.boxplot(x = 'Destination' , y = 'Price' , data = df )  
plt.xticks(rotation = 65)  
plt.show()
```




```
In [35]: df_blr_del = df[df['Route']=='BLR → DEL']
```

Lets analyse flight fare price for banglore to delhi with different airline.

```
In [36]: print(df_blr_del.groupby('Airline')['Price'].mean().sort_values(ascending = False)  
df_blr_del.groupby('Airline')['Price'].mean().sort_values(ascending = False).plot
```

```
Airline  
Vistara Premium economy    8881.000000  
Air India                  6716.757962  
Jet Airways               6498.803150  
Vistara                   5960.674286  
IndiGo                    5023.526427  
GoAir                     4767.033708  
Air Asia                  4574.280899  
SpiceJet                  4289.847059  
Name: Price, dtype: float64
```

```
Out[36]: <AxesSubplot:xlabel='Airline'>
```




Vistara premium economy airline fare price is quite high comparing with other airlines.
SpiceJet fare is very cheap.

BLG->DEL average fare is almost 6500+

```
In [37]: sns.scatterplot(x = 'total_hours' , y = 'Price' , data = df)
```


```
Out[37]: <AxesSubplot:xlabel='total_hours', ylabel='Price'>
```



Outliers


```
In [38]: df['Price'].plot(kind = 'kde')
```

```
Out[38]: <AxesSubplot:ylabel='Density'>
```



```
In [39]: sns.boxplot(x = 'Price' , data = df)
```

```
Out[39]: <AxesSubplot:xlabel='Price'>
```



```
In [40]: df.describe()
```

```
Out[40]:
```

	Total_Stops	Price	dep_hour	dep_min	arrival_hour	arrival_min	Durati
count	10461.000000	10461.000000	10461.000000	10461.000000	10461.000000	10461.000000	10461.000000
mean	0.802027	9025.962527	12.479208	24.402543	13.387917	24.720390	10.479208
std	0.659900	4624.295514	5.727034	18.814954	6.855835	16.571178	8.814954
min	0.000000	1759.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	5224.000000	8.000000	5.000000	8.000000	10.000000	5.000000
50%	1.000000	8266.000000	11.000000	25.000000	14.000000	25.000000	8.000000
75%	1.000000	12341.000000	18.000000	40.000000	19.000000	35.000000	18.000000
max	3.000000	79512.000000	23.000000	55.000000	23.000000	55.000000	40.000000

```
In [41]: df['Price'] = np.where(df['Price'] > 40000 , df['Price'].median() , df['Price'])
```

```
In [42]: df['month'] = df['Date_of_Journey'].dt.month
df['date'] = df['Date_of_Journey'].dt.day
```

```
In [43]: cols = ['Date_of_Journey' , 'Route' , 'Dep_Time' , 'Arrival_Time' , 'Duration' ,
final_df = df.drop(columns=cols , axis = True )
final_df.head()
```

Out[43]:

	Airline	Source	Destination	Total_Stops	Price	dep_hour	dep_min	arrival_hour	arrival_m
0	IndiGo	Banglore	Delhi	0	3897.0	22	20	1	
1	Air India	Kolkata	Banglore	2	7662.0	5	50	13	
2	Jet Airways	Delhi	Cochin	2	13882.0	9	25	4	
3	IndiGo	Kolkata	Banglore	1	6218.0	18	5	23	
4	IndiGo	Banglore	Delhi	1	13302.0	16	50	21	

```
In [44]: final_df = pd.get_dummies(final_df , drop_first=True)
```

```
In [45]: final_df.head()
```

Out[45]:

	Total_Stops	Price	dep_hour	dep_min	arrival_hour	arrival_min	total_hours	month	date	A
0	0	3897.0	22	20	1	10	2.833333	3	24	
1	2	7662.0	5	50	13	15	7.416667	1	5	
2	2	13882.0	9	25	4	25	19.000000	9	6	
3	1	6218.0	18	5	23	30	5.416667	12	5	
4	1	13302.0	16	50	21	35	4.750000	1	3	

5 rows × 28 columns


Correlation


```
In [46]: final_df.head()
```


Out[46]:

	Total_Stops	Price	dep_hour	dep_min	arrival_hour	arrival_min	total_hours	month	date	A
0	0	3897.0	22	20	1	10	2.833333	3	24	
1	2	7662.0	5	50	13	15	7.416667	1	5	
2	2	13882.0	9	25	4	25	19.000000	9	6	
3	1	6218.0	18	5	23	30	5.416667	12	5	
4	1	13302.0	16	50	21	35	4.750000	1	3	

5 rows × 28 columns



```
In [47]: plt.figure(figsize= (20,12))
sns.heatmap(final_df.corr() , annot = True , cmap = 'coolwarm' , vmin = -1 , vmax = 1)
```



```
In [48]: final_df.corr()['Price'].sort_values(ascending = False )[1:]
```

```
Out[48]: Total_Stops                0.627674
total_hours                0.532665
Airline_Jet Airways        0.433802
Destination_Cochin         0.279750
Source_Delhi               0.279750
Airline_Multiple carriers  0.156851
Airline_Air India          0.056779
arrival_hour               0.038111
Source_Kolkata             0.022219
month                      0.020392
Airline_Multiple carriers Premium economy 0.019542
dep_hour                   0.005497
Airline_Vistara Premium economy -0.000083
Airline_Jet Airways Business -0.003913
Airline_Trujet            -0.010776
dep_min                   -0.032893
Airline_Vistara           -0.058872
arrival_min               -0.090399
Airline_GoAir             -0.097666
Destination_Delhi         -0.131812
date                      -0.172852
Source_Chennai            -0.185513
Destination_Kolkata       -0.185513
Source_Mumbai             -0.238543
Destination_Hyderabad     -0.238543
Airline_SpiceJet          -0.307391
Airline_IndiGo            -0.371606
Name: Price, dtype: float64
```

```
In [ ]:
```

```
In [49]: # dropping price columns
X = final_df.drop('Price' , 1 )
# standarization of price columns
y = np.log1p(final_df['Price'])
```

```
In [50]: #train test split with 33% test values
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_
```

```
In [51]: X_train.shape , X_test.shape
```

```
Out[51]: ((7008, 27), (3453, 27))
```

```
In [52]: from sklearn.linear_model import LinearRegression,Ridge,Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor,GradientBoostingRegressor,AdaB
from sklearn.metrics import r2_score , mean_absolute_error , mean_absolute_perce
```

```
In [53]: def model_feature(model):  
    model.fit(X_train , y_train)  
    y_pred = model.predict(X_test)  
    print(str(model)[0 : -2] + ' ' 'Model')  
    print('r2_score:{}'.format(round(r2_score(y_test , y_pred) , 2)))  
    print('MAE',round(mean_absolute_error(y_test , y_pred) , 2))  
    print('MAPE' , round(mean_absolute_percentage_error(y_test , y_pred) , 2))  
    print('MSE' , round(mean_squared_error(y_test , y_pred) , 2))
```

LinearRegression

```
In [54]: model_feature(LinearRegression())
```

```
LinearRegression Model  
r2_score:0.69  
MAE 0.21  
MAPE 0.02  
MSE 0.08
```

Lasso

```
In [55]: model_feature(Lasso())
```

```
Lasso Model  
r2_score:0.31  
MAE 0.34  
MAPE 0.04  
MSE 0.18
```

Ridge

```
In [56]: model_feature(Ridge())
```

```
Ridge Model  
r2_score:0.69  
MAE 0.21  
MAPE 0.02  
MSE 0.08
```

KNeighborsRegressor

```
In [57]: model_feature(KNeighborsRegressor())
```

```
KNeighborsRegressor Model  
r2_score:0.7  
MAE 0.21  
MAPE 0.02  
MSE 0.08
```

RandomForestRegressor

```
In [58]: model_feature(RandomForestRegressor())
```

```
RandomForestRegressor Model  
r2_score:0.86  
MAE 0.13  
MAPE 0.01  
MSE 0.04
```

GradientBoostingRegressor

```
In [59]: model_feature(GradientBoostingRegressor())
```

```
GradientBoostingRegressor Model  
r2_score:0.83  
MAE 0.16  
MAPE 0.02  
MSE 0.04
```

AdaBoostRegressor

```
In [60]: model_feature(AdaBoostRegressor())
```

```
AdaBoostRegressor Model  
r2_score:0.71  
MAE 0.22  
MAPE 0.03  
MSE 0.07
```

ExtraTreesRegressor

```
In [61]: model_feature(ExtraTreesRegressor())
```

```
ExtraTreesRegressor Model  
r2_score:0.82  
MAE 0.14  
MAPE 0.02  
MSE 0.05
```

VotingRegressor

```
In [62]: from sklearn.ensemble import VotingRegressor, StackingRegressor

rf = RandomForestRegressor(n_estimators=350, random_state=3, max_samples=0.5, max_features=0.5)
gbdt = GradientBoostingRegressor(n_estimators=100, max_features=0.5)
# xgb = XGBRegressor(n_estimators=25, learning_rate=0.3, max_depth=5)
et = ExtraTreesRegressor(n_estimators=100, random_state=3, max_samples=0.5, max_features=0.5)

model_vr = VotingRegressor([('rf', rf), ('gbdt', gbdt), ('et', et)], weights=[5, 1, 1])

model_vr.fit(X_train, y_train)

y_pred = model_vr.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.8695471971376612

MAE 0.13300083528164577

StackingRegressor

```
In [63]: from sklearn.ensemble import VotingRegressor, StackingRegressor

estimators = [
    ('rf', RandomForestRegressor(n_estimators=350, random_state=3, max_samples=0.5, max_features=0.5)),
    ('gbdt', GradientBoostingRegressor(n_estimators=100, max_features=0.5)),
]

pipe = StackingRegressor(estimators=estimators, final_estimator=Ridge(alpha=100))

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.866720440181483

MAE 0.13688271524270923

RandomForestRegressor

```
In [64]: model_rfr = RandomForestRegressor()
model_rfr.fit(X_train , y_train)
model_rfr.predict(X_test)
r2_score(y_test,y_pred)
```

Out[64]: 0.866720440181483

```
In [65]: model_list = [LinearRegression() , Ridge() , Lasso() , KNeighborsRegressor() , DecisionTreeRegressor()]
model_list1 = []
R2_score = []
mae = []
mape = []
mse = []

for model in model_list:
    model_list1.append(str(model)[0:-2])
    model.fit(X_train , y_train)
    y_pred = model.predict(X_test)
    R2_score.append(round(r2_score(y_test , y_pred) , 2))
    mae.append(round(mean_absolute_error(y_test , y_pred) , 2))
    mape.append(round(mean_absolute_percentage_error(y_test , y_pred) , 2))
    mse.append(round(mean_squared_error(y_test , y_pred) , 2))
```

```
In [66]: dict = {'Model':model_list1, 'R2_score':R2_score , 'MAPE':mape , 'MAE':mae , 'MSE':mse}
model_df = pd.DataFrame(dict).sort_values(ascending = False , by = 'R2_score')
model_df
```

Out[66]:

	Model	R2_score	MAPE	MAE	MSE
5	RandomForestRegressor	0.86	0.01	0.13	0.04
6	GradientBoostingRegressor	0.83	0.02	0.16	0.04
8	ExtraTreesRegressor	0.82	0.02	0.14	0.05
4	DecisionTreeRegressor	0.77	0.02	0.15	0.06
7	AdaBoostRegressor	0.71	0.03	0.23	0.08
3	KNeighborsRegressor	0.70	0.02	0.21	0.08
0	LinearRegression	0.69	0.02	0.21	0.08
1	Ridge	0.69	0.02	0.21	0.08
2	Lasso	0.31	0.04	0.34	0.18

we are getting Best R2_score with RandomForestRegressor

Exporting the Model

```
In [67]: import pickle

pickle.dump(final_df,open('final_df.pkl','wb'))
pickle.dump(model_rfr,open('model.pkl','wb'))
```

```
In [ ]:
```